
Beyond the 7.8% Lean Baseline: Extending AlgoVeri with Frontier Models, Coding Agents, and Evaluator Audits

Joey Li¹

Abstract

AlgoVeri established a challenging baseline for verified code generation on 77 aligned classical algorithms: its strongest reported Lean system achieved 9.1% compiler verification and 7.8% full marks after semantic filtering. We extend that study along three axes: newer frontier models, a compiler-using coding agent, and an audit of the evaluation pipeline. Complete-output direct conditions now compile 92.2–98.7% of the Lean tasks, but only 31.2–57.1% receive full marks, depending on generator and semantic judge. Thus, the dominant bottleneck shifts from constructing a compiling Lean artifact to preserving the requested algorithm under an extensional formal contract. On a controlled 65-task subset without teacher-provided proof holes, a constrained Claude Code + Opus 5 agent compiles 63/65 tasks and passes 53/65 under GPT-5.4 and 55/65 under GPT-5.6-sol at temperature zero, compared with 39/65 and 44/65 for the strongest complete direct condition.

These results confirm AlgoVeri’s algorithmic-fidelity gap, extend its finding that frontier systems exploit iterative repair, and revise the characterization of Lean’s barrier: compilation is no longer the principal failure mode for the evaluated frontier systems, while semantic degeneracy remains pervasive. We additionally find that one immutable scaffold is malformed, 12 tasks expose a contradiction between teacher-permitted `sorry` and the judge policy, and two temperature-zero judges disagree on 22/334 compiled condition–task pairs. We release all generations, judge views, case-level analyses, an agent harness, and an interactive audit dashboard. Our evidence motivates a multi-axis protocol that re-

ports output coverage, formal verification, algorithmic fidelity, and evaluator robustness separately.

1. Introduction

Formal verification offers a stronger correctness signal than unit tests: a verifier checks a program against a mathematical contract. Yet a verified program is only as informative as that contract. A generated Lean file can type-check while implementing the wrong algorithm, delegating to a library theorem, invoking a nonconstructive specification oracle, or retaining an unverified teacher placeholder. Consequently, verified code generation must measure both *formal correctness* and *algorithmic fidelity*.

AlgoVeri introduced an unusually controlled setting for this question: 77 textbook algorithms with aligned natural-language requirements and functional contracts in Dafny, Verus, and Lean (Zhao et al., 2026). The original evaluation established three findings that motivate this work. First, vericoding remained far from solved, with the strongest Lean condition obtaining only 9.09% compiler verification and 7.79% full marks. Second, compiler verification systematically exceeded semantic acceptance, revealing an algorithmic-fidelity gap. Third, frontier models converted iterative verifier feedback into progress, whereas weaker systems saturated early; Lean failures combined proof search and logical reasoning barriers.

We ask what remains true after substituting newer frontier models and a modern coding agent into the same Lean task suite. The answer is not simply that scores increase. Compilation rises above 92% for every complete direct condition, yet semantic success lags by 42–66 percentage points. The capability frontier has therefore moved across the compiler boundary without eliminating the benchmark’s central fidelity problem. A constrained coding agent closes much of this gap, but its remaining failures reveal a stable optimization pressure toward brute-force solvers, classical choice, and verified fallbacks. Meanwhile, evaluator defects introduce measurement error precisely where scores are becoming large enough for reliable ranking to matter.

¹Independent Researcher. Correspondence to: Joey Li <joeyli.shsid@gmail.com>.

This paper makes five contributions:

1. We extend AlgoVeri’s Lean evaluation with four direct frontier-model conditions and three semantic-judge conditions, placing every result against the original paper’s capability, fidelity, and repair findings.
2. We implement and evaluate an isolated Claude Code condition with constrained compiler access, adaptive repair, source ownership enforcement, and final independent verification.
3. We show a bottleneck shift: current frontier models largely overcome Lean compilation, while algorithm identity becomes the dominant source of failure.
4. We audit every compiler failure, every coding-agent semantic rejection, every judge disagreement, and all task-level outcomes. The complete matrix appears in Appendix 6; machine-readable analyses are released with the artifact(Li, 2026).
5. We identify a malformed scaffold, an ownership contradiction involving teacher-permitted `sorry`, an asymmetric repair prompt, and measurable judge dependence, then propose concrete benchmark fixes.

This is a Lean-only extension, not a reproduction of AlgoVeri’s cross-language ranking. Because model endpoints, inference controls, and semantic judges differ, comparisons to the original reported rates are descriptive historical comparisons rather than a causal estimate of model progress.

2. Background and Related Work

2.1. Lean and verified code generation

Lean 4 combines dependent type theory, an extensible theorem prover, and a functional programming language(de Moura & Ullrich, 2021). Mathlib supplies a large, community-built library of definitions, theorems, and tactics(The mathlib Community, 2020). This ecosystem enables machine-checkable proofs but also creates evaluation choices: using Mathlib is necessary for the supplied tasks, whereas calling an existing implementation of the target algorithm may violate a “from scratch” requirement. The benchmark therefore needs a semantic policy in addition to compiler verification.

Formal-reasoning benchmarks such as miniF2F(Zheng et al., 2022) and LeanDojo(Yang et al., 2023) primarily evaluate theorem proving. Verified program generation additionally requires executable definitions, functional contracts, termination, and algorithm-specific invariants. DafnyBench(Loughridge et al., 2024) and the multi-language vericoding benchmark(Bursuc et al., 2025)

demonstrate the growing interest in this broader setting. AlgoVeri contributes aligned classical algorithm tasks across verification paradigms(Zhao et al., 2026).

2.2. AlgoVeri’s two-stage Lean metric

Each Lean task contains a natural-language description and a scaffold with four model-editable marker regions: `auxcode`, `code`, `lemma`, and `proof`. The response parser copies only those regions into the teacher scaffold. Generated uses of `sorry`, `admit`, `axioms`, `unsafe definitions`, and related bypasses are rejected. A local Lean compiler then checks the merged file.

Compilation alone cannot enforce the named algorithm. AlgoVeri therefore applies an LLM semantic filter to the natural-language task and complete merged file. The judge checks whether the intended algorithm was implemented from scratch and whether the answer uses a compiler trick. This is an instance of model-based evaluation, whose scalability comes with known sensitivity to judge and prompt(Zheng et al., 2023). We retain the benchmark prompt while explicitly measuring that sensitivity.

3. Extending AlgoVeri’s Findings

Table 1 states the relationship between the original conclusions and our extension. We treat the published AlgoVeri rates as a historical reference, retain its compile-then-semantic evaluation logic, and introduce controlled analyses where the original protocol leaves evaluator uncertainty unmeasured.

We organize the empirical study around five research questions:

- RQ1** How far has the Lean capability frontier moved relative to AlgoVeri’s original 9.09% verified and 7.79% full-mark reference point?
- RQ2** Does AlgoVeri’s algorithmic-fidelity gap persist after compilation becomes common rather than exceptional?
- RQ3** Does a constrained compiler-using coding agent improve end-to-end semantic performance relative to direct API generation?
- RQ4** How stable are semantic conclusions under temperature and judge-model changes?
- RQ5** Which task, prompt, and evaluator defects threaten the validity of the measured score, and how should the protocol change?

Table 1. How this study extends the principal findings of AlgoVeri. Numerical comparisons across papers are descriptive because models and endpoint controls differ.

Original finding	Evidence in this extension	Updated interpretation
Lean is the hardest target; best full mark is 7.79%.	Complete direct conditions compile 92.2–98.7% and obtain 31.2–57.1% full marks.	The evaluated frontier has crossed much of Lean’s compilation barrier, but vericoding remains unsolved.
Compiler verification overstates algorithmic correctness.	The direct-model compiler-to-semantic gap is 41.6–66.2 percentage points on full-scope conditions.	The fidelity gap grows more visible as formal compilation becomes easier.
Frontier models benefit from iterative repair; weaker models saturate early.	All complete direct conditions continue gaining compilation successes after the first round; the agent moves from 15 to 62 successes by its fifth Lean check.	Tool-mediated repair remains valuable, but compiler feedback optimizes formal acceptance rather than algorithm identity.
Lean combines search and reasoning barriers.	Final direct compilation failures are few, while semantic failures use wrong algorithms, oracles, and fallbacks.	The bottleneck shifts from proof search toward specification alignment and algorithmic fidelity for newer frontier systems.
Expert curation and semantic filtering protect benchmark validity.	We find one uncompileable scaffold, 12 ownership-confounded tasks, and 22/334 cross-judge flips.	Task curation must be complemented by executable preflight and evaluator calibration.

4. Experimental Design

4.1. Relationship to the original protocol

Both studies use the same 77 Lean task identities, natural-language algorithm requirements, four editable regions, iterative compiler feedback, and a semantic filter after successful compilation. The direct conditions retain a maximum of 15 attempts, matching the original study’s main repair depth. Our extension differs in four ways: it is Lean-only; it evaluates later proprietary model endpoints; it compiles through a pinned local Lake environment rather than treating the container runtime as essential; and it adds an agentic treatment plus three judge configurations. We therefore compare the *shape* of the findings and report original numbers as context, but reserve paired statistical tests for conditions evaluated on the same tasks in this artifact.

4.2. Task scopes

The full scope contains all 77 Lean tasks in the public AlgoVeri repository (Zhao & contributors, 2026). Four direct conditions target the full scope. The adaptive-thinking Opus run produced only 49 saved outputs; missing outputs remain failures under the end-to-end denominator and are also reported conditionally.

The enhanced coding-agent experiment uses a deterministic source-aware filter that excludes the 12 tasks containing teacher-owned `sorry` outside all four editable regions. Its scope is therefore 65 tasks. We evaluate every generator on this same 65-task set for fair comparisons. No generation result is selected or discarded based on its later semantic verdict.

4.3. Direct generation conditions

We evaluate GPT-5.5, GPT-5.6-sol, and Claude Opus 5 through the same AlgoVeri chat-and-compiler-revision loop. Each task receives one independent pass, up to 15 generation or repair attempts, and compiler error feedback after an invalid attempt. The OpenAI conditions use medium reasoning and a 32,768-token maximum output. Opus is evaluated both with adaptive thinking and with thinking disabled. Provider-specific reasoning settings are not assumed to be equivalent.

The initial prompt explains the algorithm requirement and prohibits compiler bypasses. After a failed attempt, the stock Lean revision prompt asks for a complete standalone file that addresses the compiler errors. Section 7.3 discusses the semantic information lost in this shorter revision instruction.

4.4. Enhanced coding-agent condition

The coding-agent condition uses Claude Code 2.1.220 with Opus 5, medium effort, and an isolated workspace per task. The agent may read task files, edit only `Solution.lean`, and invoke only a constrained `./check.sh`; network, Git, package installation, and arbitrary shell commands are blocked. Candidate content is copied only from the four editable regions into the original scaffold. The harness then rejects prohibited terms and independently recompiles the merged result.

The enhanced protocol provides an initial budget of USD 2. If independent compilation fails, one fresh compiler-repair session receives the exact verifier feedback and up to another USD 2 while preserving the workspace. The prompt explicitly prohibits alternative algorithms, specifi-

cation oracles, brute-force enumeration, sorting as a substitute, `Nat.find`, `Classical.choose`, and `decide` over the target property. It also requires an early Lean check and a final requirement-by-requirement semantic audit. The final saved result is always independently checked; the agent’s claim of success is not trusted.

Before this enhanced run, two targeted pilot batches covered 10 and 20 cases that prior direct conditions had failed. Those pilots are reported separately because their selection and prompt differ from the final 65-task condition.

4.5. Lean environment

All candidates are compiled locally without Apptainer. The runner places a temporary file inside a pinned Lake project and invokes `lake env lean`. The environment uses Lean 4.25.0-rc2 (commit 744f9806), Lake 5.0.0, and Mathlib tag v4.25.0-rc2 (commit 766e19e7). The host runs Ubuntu 22.04.5 LTS on an AMD Ryzen Threadripper PRO 7985WX (64 physical cores, 128 threads) with 502 GiB RAM. The visible display adapter is AMD/ATI; generation and semantic inference use remote APIs, so local GPU performance is not part of the treatment.

4.6. Semantic judges

Every saved generation output is evaluated with the same AlgoVeri semantic prompt under three conditions:

1. GPT-5.4, temperature 1 (the initial rerun configuration);
2. GPT-5.4, temperature 0;
3. GPT-5.6-sol, temperature 0 and reasoning effort `none`.

GPT-5.6-sol rejects an explicit temperature when reasoning is enabled in this API, so `none` is required to isolate temperature zero. The temperature-zero outputs are stored in distinct directories and do not overwrite the initial judgments. Up to five turns are allowed only to repair malformed `<analysis>` and `<conclusion>` formatting; this is not a vote across independent samples.

4.7. Metrics and analysis

We report four mutually exclusive observable states: missing output, saved output that does not compile, compiled output rejected by the semantic judge, and compiled output accepted by the judge. “Semantic pass” always means the conjunction of compiler success, a parsed judgment, and a `YES` verdict. Rates use the stated native or common scope as denominator; we additionally show output-conditional rates where missingness is material.

Judge agreement is measured only on compiled outputs

evaluated by both judges. We report raw agreement and Cohen’s κ (Cohen, 1960). Paired agent/direct comparisons on the same 65 tasks use the exact two-sided McNemar test (McNemar, 1947); these descriptive p -values do not make the fixed task suite a random population. Selected binomial intervals are 95% Wilson intervals (Wilson, 1927). All aggregates are regenerated from JSON by `scripts/analyze_rerun_results.py`; no table was transcribed from terminal logs.

5. Results: From Proof Search to Fidelity

5.1. The Lean capability frontier has moved

Table 2 answers RQ1. AlgoVeri’s strongest original Lean result was Gemini-3 Flash at 7/77 compiler-verified and 6/77 full-mark tasks (9.09% and 7.79%) (Zhao et al., 2026). In our later evaluation, when an answer is returned, compilation is high: GPT-5.5 compiles 75/77 (97.4%; Wilson 95% CI [91.0%, 99.3%]), GPT-5.6-sol 71/77 (92.2%), and no-thinking Opus 76/77 (98.7%). The protocols are not identical, so these numbers do not isolate model improvement; they nevertheless show that the observed failure regime has changed. All 49 saved adaptive-thinking Opus outputs compile, but 28 tasks have no output, reducing end-to-end compilation to 63.6%. In its final one-attempt sweep over the missing set, 24 requests exhausted 32,768 output tokens with thinking blocks only and four ended in provider overload.

Disabling Opus thinking eliminates this coverage failure and is not offset by worse quality on the matched 49 returned tasks. On those tasks, thinking versus no-thinking semantic passes are 29 versus 30 (5.4-T1), 29 versus 31 (5.4-T0), and 31 versus 35 (5.6-T0). Thus this run provides no evidence that adaptive thinking improves the matched outcome, while it substantially reduces completion reliability.

5.2. The fidelity gap widens after compilation

The central result for RQ2 is the size of the compilation-semantic gap. AlgoVeri introduced this gap as evidence that formal verification alone does not guarantee the requested algorithm. Our extension shows that the issue becomes more pronounced, not less, once models compile reliably. For example, GPT-5.5 compiles 75 tasks but all three judges accept only 24. No-thinking Opus compiles 76, while acceptance ranges from 36 to 44. Even conditional on compilation, GPT-5.4 at temperature zero accepts 32.0% of GPT-5.5, 38.0% of GPT-5.6-sol, and 51.3% of no-thinking Opus outputs. The compiler verifies the submitted theorem; it does not establish that the program realizes the natural-language algorithm.

The 65-task scope removes the known teacher-sorry con-

Table 2. Native-scope results. Semantic columns show pass counts for GPT-5.4 at temperature 1 (5.4-T1), GPT-5.4 at temperature 0 (5.4-T0), and GPT-5.6-sol at temperature 0 (5.6-T0). Percentages use the native scope, not only returned outputs.

Generation condition	Output	Lean compile	5.4-T1	5.4-T0	5.6-T0
GPT-5.5	77/77	75/77 (97.4%)	24/77 (31.2%)	24/77 (31.2%)	24/77 (31.2%)
GPT-5.6-sol	77/77	71/77 (92.2%)	25/77 (32.5%)	27/77 (35.1%)	30/77 (39.0%)
Opus 5, thinking	49/77	49/77 (63.6%)	29/77 (37.7%)	29/77 (37.7%)	31/77 (40.3%)
Opus 5, no thinking	77/77	76/77 (98.7%)	36/77 (46.8%)	39/77 (50.6%)	44/77 (57.1%)
Claude Code + Opus 5	65/65	63/65 (96.9%)	53/65 (81.5%)	53/65 (81.5%)	55/65 (84.6%)

found and makes all final conditions directly comparable (Table 3). GPT-5.5 compiles all 65 yet passes only 24. No-thinking Opus is the strongest direct condition at 39/65 under 5.4-T0 and 44/65 under 5.6-T0.

5.3. Coding-agent gains

The enhanced coding agent answers RQ3. It compiles 63/65 and passes 53/65 under both GPT-5.4 conditions and 55/65 under GPT-5.6-sol. Its 5.6-T0 semantic rate is 84.6% (Wilson 95% CI [73.9%, 91.4%]), versus 67.7% ([55.6%, 77.8%]) for no-thinking Opus on the same tasks. The agent has 13 unique passes against two unique passes for no-thinking Opus under 5.6-T0 (exact McNemar $p = 0.0074$). The corresponding discordances are 18–1 under 5.4-T1 ($p = 7.6 \times 10^{-5}$) and 16–2 under 5.4-T0 ($p = 0.0013$).

This is an agentic condition, not a like-for-like model replacement. It receives tool access, multiple turns, persistent files, explicit semantic reminders, and an adaptive budget. The observed gain therefore estimates the whole harness treatment. The final 65-task run costs USD 70.26 in recorded Claude Code charges (mean USD 1.08 per task) and contains 6.24 summed agent-hours; direct API dollar cost and duration were not recorded and cannot be compared.

The earlier targeted pilots compiled 4/10 and 13/20, respectively; combined, 17/30 compiled and 14/30 passed the 5.4-T1 judge at a recorded cost of USD 41.48. After adding early checks, a semantic self-audit, a preserved-workspace repair pass, exact verifier feedback, and final independent verification, the broader enhanced condition reaches 63/65 compilation and 53/65 5.4-T1 semantic success. Because task selection and several settings changed, this is an engineering progression rather than an ablation.

5.4. Frontier repair remains effective but misaligned

Iterative compiler feedback matters: only 19–24 of 77 full-output direct conditions compile on the first attempt, rising to 63–69 by attempt five. GPT-5.5 gains no further success after attempt ten, and GPT-5.6-sol gains one. The coding agent reaches 62/65 by its fifth observed Lean check and one more by its tenth. These are cumulative success-at-repair curves, not pass@ k estimates from independent

samples.

This pattern extends AlgoVeri’s distinction between frontier repair and early saturation. The evaluated frontier systems continue converting Lean feedback into compilation gains, and the coding-agent workspace makes that feedback operationally effective. However, compiler feedback contains no signal about whether a different algorithm or guarded fallback violates the natural-language request. Repair is thus effective for the formal objective while remaining partially misaligned with the full- mark objective.

5.5. Judge robustness

Table 5 answers RQ4. Temperature zero does not eliminate evaluator variation. Across all five generation conditions, GPT-5.4 at temperatures 1 and 0 flips 13 of 334 compiled pairs. The two temperature-zero judges flip 22 pairs. GPT-5.6-sol is more permissive in aggregate: 17 GPT-5.6-only passes versus five GPT-5.4-only passes, a net gain of 12.

Per-condition agreement ranges from 90.1% to 96.8% for the two zero- temperature judges. Exact disagreement cases and directional flips appear in Appendix 7. Several are substantively ambiguous. For example, the agent’s `bubble_sort` is accepted at 5.4-T1, rejected at 5.4-T0 as selection-sort-like, and accepted at 5.6-T0. Its `polymul_naive` is rejected by GPT-5.4 for an unused placeholder definition but accepted by GPT-5.6-sol based on the actual convolution implementation. A single uncanceled LLM judge should therefore not be treated as ground truth.

5.6. Task-level difficulty

Sixteen tasks are stable successes across every generation condition, every available output, and all three judges: `bracket_matching`, `bst_insert`, `bst_search`, `bst_zig`, `bst_zigzag`, `bst_zigzig`, `integer_exponential`, `knapsack_01`, `llrbt_flipcolor`, `llrbt_rotateleft`, `matrix_multiplication`, `queue_dequeue`, `queue_enqueue`, `ringbuffer_enqueue`, `stack_pop`, and `stack_push`.

At the other extreme, under 5.6-T0 no

Table 3. Common 65-task scope with no teacher-owned `sorry`. This table should be used for comparisons with the enhanced coding agent.

Generation condition	Output	Lean compile	5.4-T1	5.4-T0	5.6-T0
GPT-5.5	65	65	24 (36.9%)	24 (36.9%)	24 (36.9%)
GPT-5.6-sol	65	60	25 (38.5%)	27 (41.5%)	29 (44.6%)
Opus 5, thinking	47	47	29 (44.6%)	29 (44.6%)	31 (47.7%)
Opus 5, no thinking	65	65	36 (55.4%)	39 (60.0%)	44 (67.7%)
Claude Code + Opus 5	65	63	53 (81.5%)	53 (81.5%)	55 (84.6%)

Table 4. Cumulative compilation success by attempt/check count. Direct runs count LLM generation attempts; Claude Code counts observed `./check.sh` invocations, so the curves are descriptive rather than directly cost-equivalent.

Condition	1	2	3	5	10	15
GPT-5.5	21	43	57	67	75	75
GPT-5.6-sol	19	37	53	63	70	71
Opus, thinking	33	45	49	49	49	49
Opus, no thinking	24	43	58	69	75	76
Claude Code	15	37	45	62	63	63

condition passes eight of the 65 controlled tasks: `ac_automata`, `dijkstra`, `edmond_karp`, `kruskal`, `llrbt_delete`, `longest_palindrome_substring`, `push_relabel`, and `scc_tarjan`. These concentrate graph invariants, balanced-tree preservation, and algorithm-identity requirements that are easy to replace with extensional specification solvers.

6. Failure Analysis

6.1. Compilation failures

The small number of final compilation failures fall into four distinct classes.

1. **Malformed immutable scaffold.** `trie_search` fails identically for GPT-5.5, GPT-5.6-sol, and no-thinking Opus because the teacher file itself contains a duplicate declaration and undefined identifier (Section 7.1).
2. **Response-schema failure.** GPT-5.6-sol’s final `bellman_ford`, `kruskal`, and `llrbt_delete` responses lack the required code/proof marker sections. The parser therefore has no candidate to compile.
3. **Unresolved Lean proof errors.** GPT-5.5’s `trie_insert` ends with tactic and recursion-depth errors. GPT-5.6-sol’s `llrbt_insert` leaves type mismatches and unsolved goals, while `prim` references an undefined helper theorem.
4. **Prohibited generated placeholder.** The enhanced agent’s two failures, `dijkstra` and `kruskal`, reach final candidate validation with an editable proof section still containing `sorry`. The harness correctly records

these attempts as failures rather than silently omitting files.

The common compilation problem is therefore not lack of machine resources. It is a mixture of benchmark validity, output-protocol compliance, and unfinished proof construction. More repair rounds cannot solve immutable scaffolding, and schema-aware constrained decoding may be more useful than additional proof search for missing marker regions.

6.2. Why compiler-accepted programs fail semantically

Reviewing the judge analyses and generated Lean reveals five recurring patterns.

Different algorithm, correct extensional behavior.

Many answers prove the supplied input–output property with a simpler algorithm: insertion or selection behavior for a named sort, direct convolution for Karatsuba, naive substring search for KMP, mutual reachability for Tarjan SCC, or exhaustive shortest paths instead of Dijkstra. Lean correctly verifies the theorem, but the answer does not satisfy the natural-language algorithm constraint.

Existing verified primitive. Examples include computing GCD with `Nat.gcd` and marking primes with `decide (Nat.Prime i)` rather than implementing trial division or a sieve. These are mathematically valid terms and useful Lean programming patterns, but they violate the benchmark’s from-scratch policy.

Brute force or nonconstructive specification oracle.

Several answers enumerate all binary vectors for Gaussian elimination, all edge subsets for matching, or all flow tables

Table 5. Pooled semantic-judge agreement on compiled outputs. A generation of the same benchmark task is a separate pair; pooled observations are not assumed independent.

Comparison	n	Agree	Rate	κ	Left only	Right only
5.4-T1 vs 5.4-T0	334	321	96.1%	0.922	4	9
5.4-T0 vs 5.6-T0	334	312	93.4%	0.868	5	17

for maximum flow. Others use `Classical.choose` or finite search to obtain an object known to satisfy the post-condition. This behavior exposes contracts that specify *what* result is correct but cannot encode *which* algorithm computed it.

Guarded fallback. The enhanced agent often implements a plausible target algorithm but protects it with a verified reference solver. The exported function returns the candidate only if a certificate succeeds; otherwise it rebuilds or brute-forces a correct answer. Such fallbacks maximize theorem completion while transferring correctness away from the named algorithm.

Placeholder or weak helper. An unused helper defined as 0 or `True`, or a proof that exploits a weak/vacuous precondition, may compile without providing the intended development. The judges differ on whether an unused placeholder alone warrants rejection, contributing to the disagreement in Section 5.

6.3. Residual coding-agent failures

The two temperature-zero judges agree on eight semantic rejections among the 63 compiled agent outputs:

- `ac_automata`: dynamic suffix search rather than an explicit trie with BFS-computed failure links;
- `edmond_karp`: a nonconstructive maximum-flow fallback using classical choice;
- `llrbt_delete`: validate the intended deletion, otherwise rebuild from a filtered inorder list;
- `llrbt_insert`: flatten, insert into a sorted list, and rebuild instead of local LLRBT insertion;
- `longest_palindrome_substring`: guard a Manacher-style candidate with exhaustive search;
- `max_matching`: enumerate all edge subsets instead of using augmenting paths;
- `push_relabel`: fall back to exhaustive enumeration of bounded flows;
- `scc_tarjan`: compute mutual-reachability classes without Tarjan’s stack, indices, or low-link values.

These failures persist despite explicit prompt prohibitions on exactly these strategies. Compiler interaction substantially improves theorem completion, but verifier feedback alone rewards contract satisfaction rather than algorithm identity. A semantic self-audit reduces the problem; it cannot provide a machine-checkable enforcement mechanism.

7. Evaluation Audit: Extending Quality Control

AlgoVeri emphasizes expert specification curation and representative formal well-formedness checks. Our audit extends that quality-control question from logical contracts to the executable evaluation pipeline: whether immutable scaffolds compile, whether authorship is preserved for semantic adjudication, and whether judge decisions are robust. RQ5 reveals three independent validity risks.

7.1. Malformed `trie_search` scaffold

The supplied file declares `is_valid_key` twice at lines 59 and 62. It then uses an undefined `well_formed` at lines 69 and 79. These locations are outside all four model-editable sections. Direct compilation of otherwise different generated answers consistently reports:

```
line 62: 'is_valid_key' has already been
        declared
line 69: Unknown identifier 'well_formed'
line 79: Unknown identifier 'well_formed'
```

The generation prompt simultaneously requires all non-editable content to remain unchanged, so no conforming answer can compile. This is a benchmark defect, not a model failure. It lowers end-to-end compilation for every complete full-scope condition by one and should be repaired or removed before model comparison.

7.2. Teacher-permitted `sorry` versus semantic rejection

The generation prompt expressly warns that some `sorry` terms occur outside the editable sections, mostly in `decreasing_by` termination proofs, and instructs the model not to change them. The response parser likewise checks prohibited terms only in model-owned regions. In contrast, the semantic prompt lists visible `sorry` as a cheating signal and receives the complete merged file with-

out machine-readable ownership annotations.

Twelve tasks have teacher-owned `sorry`: `segmenttree_build`, `segmenttree_modify`, `segmenttree_query`, `splaytree_splay`, the three ternary-search-tree tasks, the three trie tasks, `unionfind_find`, and `unionfind_linkroots`. Across the four direct conditions, 34 returned answers on this subset compile. GPT-5.4 assigns zero semantic passes at both temperatures. Its analysis of GPT-5.6-sol’s `ternarysearchtree_delete`, for example, says the delete and eager pruning are correct but rejects the answer solely because teacher definitions contain termination `sorry`.

GPT-5.6-sol at temperature zero accepts one of the same 34 answers: that exact `ternarysearchtree_delete`. Its analysis correctly distinguishes the two teacher-level termination holes from the student’s complete implementation and proof. This isolated acceptance confirms that ownership-aware adjudication is possible, but the prompt does not reliably elicit it. Scores on these 12 tasks conflate student quality with teacher scaffold policy.

The clean fix is not to allow generated proof holes. It is to give the judge only the student-owned sections plus necessary read-only context, or to annotate ownership and state explicitly that pre-existing teacher holes are permitted. The enhanced 65-task condition excludes the confounded tasks instead of retroactively changing their labels.

7.3. Revision-prompt asymmetry

The initial generation prompt emphasizes the exact natural-language algorithm and prohibits bypasses. The subsequent Lean revision prompt says only to correct compiler messages, provide a complete standalone file, and retain marker structure. It does not restate algorithm identity, from-scratch implementation, or the distinction between teacher and student `sorry`. The original system message remains in chat, so this is an attentional asymmetry rather than a logical permission to cheat. Nevertheless, each repair turn makes compiler acceptance the immediate objective, while the later semantic judge applies the stricter algorithmic criterion.

The observed failures are consistent with this incentive: extensive repairs converge to classical choice, exhaustive solvers, alternate algorithms, or guarded fallbacks. The enhanced agent restates semantic constraints during both initial and repair phases and performs substantially better, although tool access and budget confound a causal attribution to prompt design alone. A controlled ablation is needed to measure the prompt effect.

7.4. Metric and reporting implications

The term “success rate of Lean” is ambiguous. This audit distinguishes at least: output coverage, parse compliance, compiler verification, absence of prohibited student terms, algorithmic fidelity, and judge agreement. A single number hides where the system failed and can reverse rankings when missing outputs or teacher-owned holes are handled differently. We recommend a structured result record rather than a binary leaderboard field.

8. Discussion and Recommendations

1. **Preflight every scaffold.** Compile a teacher-filled reference after erasing only the four student holes; reject duplicate declarations, missing identifiers, and version drift before release.
2. **Make source ownership explicit.** Preserve marker provenance through semantic evaluation. Judges should never attribute immutable teacher terms to the student.
3. **Align all repair prompts.** Restate the exact-algorithm, from-scratch, and no-oracle constraints whenever compiler feedback is supplied.
4. **Separate formal and semantic metrics.** Report output, parse, compilation, prohibited-term validation, and semantic fidelity independently.
5. **Calibrate semantic evaluation.** Use temperature-zero runs, at least two judges or blinded human adjudication for disagreements, and publish the judge model, API mode, prompt, and sampling parameters.
6. **Test algorithm identity where possible.** Add structural obligations, trace properties, complexity/resource checks, or required intermediate invariants so that a brute-force oracle cannot satisfy the same formal contract.
7. **Record every attempt.** Empty responses, quota errors, overloads, parser failures, and prohibited-term candidates should produce durable failure artifacts rather than disappearing from the denominator.
8. **Evaluate agents as a separate treatment.** Publish tool policy, compiler-call counts, budgets, final independent verification, and costs; do not merge agentic scores with direct prompting.

9. Limitations and Threats to Validity

Single samples. Each generation condition contains one pass per task. We do not estimate variance over random seeds, API nondeterminism, or temporal model updates. Temperature zero is a sampling setting, not a guarantee of bitwise determinism.

LLM semantic labels. No expert manually relabeled all 373 condition–task pairs. The judge analyses are useful evidence, not formal proofs of algorithm identity. The measured 22 cross-judge flips demonstrate residual label uncertainty.

Unequal treatments. Provider reasoning controls and token accounting are not directly comparable. The coding agent receives tools, more interaction, and a monetary budget. Its result is an upper-level system comparison, not an isolated Opus model effect.

Adaptive-thinking attrition. The 49 returned adaptive-thinking outputs are not a random subset of 77. Conditional quality on those outputs is subject to selection bias; end-to-end rates retain missing outputs in the denominator.

Benchmark and version scope. Findings apply to the audited repository snapshot, Lean 4.25.0-rc2, and the named model endpoints in July–August 2026. They do not automatically generalize to corrected scaffolds, other Lean/Mathlib releases, other languages in AlgoVeri, or future endpoint revisions.

Cost accounting. Only Claude Code reports dollar cost in the saved records. Token fields differ by provider (for example, Anthropic thinking is not stored in the same field as OpenAI reasoning), so we release raw counters but avoid a cross-provider efficiency ranking.

10. Reproducibility and Artifact Layout

The artifact repository is <https://github.com/Joey-the-Creater/AlgoVeri-Rerun>. It includes generation JSON, semantic JSON for all three judge settings, source modifications, tests, the dashboard, and this report. The principal immutable directories are:

- `results/full_three`: GPT-5.5, GPT-5.6-sol, and thinking-Opus generation;
- `results/full_opus_no_thinking`: no-thinking Opus generation;
- `results/claude_code_opus5`: enhanced 65-task agent generation;
- `matching _semantic`, `_semantic_temp0`, and `_semantic_gpt56_temp0` directories for the three judges;
- `reports/data/result_summary.json` and `reports/data/case_results.csv`: recomputed aggregate and case-level data;

- `config/dashboard_experiments.json`: the canonical experiment catalog.

From the repository root, regenerate the analytical artifacts with:

```
python3 scripts/analyze_rerun_results.py \
--json-output reports/data/result_summary.json \
--csv-output reports/data/case_results.csv \
--latex-output reports/generated/case_tables.tex
```

Compile the paper locally or in Overleaf with:

```
cd reports
latexmk -pdf algoveri_rerun_report.tex
```

The CSV contains every judge analysis in full. The PDF matrix uses compact pass/fail codes to remain readable; the underlying generated Lean and complete conversations are preserved in the result JSON.

11. Conclusion

AlgoVeri originally showed that Lean vericoding was limited by proof search and reasoning, that compiler verification exceeded algorithmic correctness, and that frontier models benefited from repair. Our extension confirms the latter two findings while revising the first for current frontier systems. Complete direct conditions now compile more than 92% of tasks, yet attain only 31–57% full marks. The principal bottleneck has shifted from reaching the Lean compiler boundary to preserving the named algorithm under contracts that admit extensionally correct substitutes.

A constrained coding agent raises full-mark performance to 82–85% on the clean 65-task scope, demonstrating that persistent tools and verifier feedback materially advance the frontier. Its residual oracle and fallback solutions also show why agentic compute alone does not solve algorithmic fidelity. Finally, scaffold, ownership, and judge-disagreement findings extend benchmark quality control beyond specification curation. Future vericoding benchmarks should treat formal verification, algorithm identity, provenance, output coverage, and evaluator robustness as distinct measurable objects.

Impact Statement

Verified-code agents may reduce the cost of producing high-assurance software, but an inflated benchmark score can encourage unsafe trust in code that proves the wrong property or implements the wrong algorithm. This study’s intended impact is improved measurement: independent com-

pilation, provenance-aware semantic review, and transparent failure artifacts. The same techniques could be used to optimize toward benchmark judges rather than genuine correctness; human review of specifications and disputed semantic labels therefore remains necessary in safety-critical deployment.

References

- Bursuc, S., Ehrenborg, T., Lin, S., Astefanoaei, L., Chiosa, I. E., Kukovec, J., Singh, A., Butterley, O., Bizid, A., Dougherty, Q., Zhao, M., Tan, M., and Tegmark, M. A benchmark for vericoding: Formally verified program synthesis. *arXiv preprint arXiv:2509.22908*, 2025. URL <https://arxiv.org/abs/2509.22908>.
- Cohen, J. A coefficient of agreement for nominal scales. *Educational and Psychological Measurement*, 20(1):37–46, 1960. doi: 10.1177/001316446002000104.
- de Moura, L. and Ullrich, S. The Lean 4 theorem prover and programming language. In *Automated Deduction—CADE 28*, pp. 625–635. Springer, 2021. doi: 10.1007/978-3-030-79876-5_37.
- Li, J. AlgoVeri-Rerun: Code, results, and reproducibility artifacts. <https://github.com/Joey-the-Creater/AlgoVeri-Rerun>, 2026.
- Loughridge, C., Sun, Q., Ahrenbach, S., Cassano, F., Sun, C., Sheng, Y., Mudide, A., Misu, M. R. H., Amin, N., and Tegmark, M. DafnyBench: A benchmark for formal software verification. *arXiv preprint arXiv:2406.08467*, 2024. URL <https://arxiv.org/abs/2406.08467>.
- McNemar, Q. Note on the sampling error of the difference between correlated proportions or percentages. *Psychometrika*, 12(2):153–157, 1947. doi: 10.1007/BF02295996.
- The mathlib Community. The Lean mathematical library. In *Proceedings of the 9th ACM SIGPLAN International Conference on Certified Programs and Proofs*, pp. 367–381. ACM, 2020. doi: 10.1145/3372885.3373824.
- Wilson, E. B. Probable inference, the law of succession, and statistical inference. *Journal of the American Statistical Association*, 22(158):209–212, 1927. doi: 10.1080/01621459.1927.10502953.
- Yang, K., Swope, A., Gu, A., Chalamala, R., Song, P., Yu, S., Godil, S., Prenger, R., and Anandkumar, A. LeanDojo: Theorem proving with retrieval-augmented language models. In *Advances in Neural Information Processing Systems*, volume 36, 2023. URL <https://arxiv.org/abs/2306.15626>.
- Zhao, H. and contributors. AlgoVeri source repository. <https://github.com/haoyuzhao123/algoveri>, 2026. Accessed 2026-08-06.
- Zhao, H., Yang, Z., Li, J., He, D., Li, Z., Jin, C., Veeravalli, V. V., Gupta, A., and Arora, S. AlgoVeri: An aligned benchmark for verified code generation on classical algorithms. *Proceedings of the 43rd International Conference on Machine Learning*, 2026. URL <https://arxiv.org/abs/2602.09464>. Accepted to ICML 2026; arXiv:2602.09464v2.
- Zheng, K., Han, J. M., and Polu, S. miniF2F: A cross-system benchmark for formal olympiad-level mathematics. In *International Conference on Learning Representations*, 2022. URL <https://arxiv.org/abs/2109.00110>.
- Zheng, L., Chiang, W.-L., Sheng, Y., Zhuang, S., Wu, Z., Zhuang, Y., Lin, Z., Li, Z., Li, D., Xing, E. P., Zhang, H., Gonzalez, J. E., and Stoica, I. Judging LLM-as-a-judge with MT-Bench and Chatbot Arena. In *Advances in Neural Information Processing Systems*, volume 36, 2023. URL <https://arxiv.org/abs/2306.05685>.

A. Native-Scope Failure Inventories

GPT-5.5 compilation failures (2). `trie_insert`, `trie_search`.

GPT-5.6-sol compilation failures (6). `bellman_ford`, `kruskal`, `llrbt_delete`, `llrbt_insert`, `prim`, `trie_search`.

No-thinking Opus compilation failures (1). `trie_search`.

Enhanced Claude Code compilation failures (2). `dijkstra`, `kruskal`.

Thinking-Opus missing outputs (28). `bellman_ford`, `bfs`, `coin_change`, `cycle_detection`, `dfs`, `dijkstra`, `jump_game`, `knapsack_unbounded`, `kruskal`, `lca`, `llrbt_delete`, `llrbt_insert`, `maxheap_popmax`, `prim`, `push_relabel`, `scc_tarjan`, `segmenttree_modify`, `segmenttree_query`, `ternarysearchtree_delete`, `ternarysearchtree_insert`, `ternarysearchtree_search`, `topological_sort`, `trial_division_optimized`, `trie_delete`, `trie_insert`, `trie_search`, `unionfind_find`, `unionfind_linkroots`.

B. Complete Case Matrix and Judge Disagreements

Table 6. Complete case-level outcome matrix. A cell is C:abc when Lean compiles; *a*, *b*, and *c* are semantic outcomes for GPT-5.4 at temperature 1, GPT-5.4 at temperature 0, and GPT-5.6-sol at temperature 0, respectively (P=pass, F=fail). X denotes a saved output that does not compile, M a missing output, and – an out-of-scope case.

Task	GPT-5.5	GPT-5.6-sol	Opus think	Opus no-think	Claude Code
ac_automata	C:FFF	C:FFF	C:FFF	C:FFF	C:FFF
bellman_ford	C:FFF	X	M	C:FPP	C:PPP
bfs	C:FFF	C:FFF	M	C:FFF	C:PPP
binary_search	C:FFF	C:PPP	C:PPP	C:PPP	C:PPP
bipartite_check	C:FFF	C:FFF	C:FPP	C:FPP	C:PPP
bracket_matching	C:PPP	C:PPP	C:PPP	C:PPP	C:PPP
bst_delete	C:PPP	C:FFF	C:PPP	C:PPP	C:PPP
bst_insert	C:PPP	C:PPP	C:PPP	C:PPP	C:PPP
bst_search	C:PPP	C:PPP	C:PPP	C:PPP	C:PPP
bst_zig	C:PPP	C:PPP	C:PPP	C:PPP	C:PPP
bst_zigzag	C:PPP	C:PPP	C:PPP	C:PPP	C:PPP
bst_zigzig	C:PPP	C:PPP	C:PPP	C:PPP	C:PPP
bubble_sort	C:FFF	C:FFF	C:PPP	C:PPP	C:FPP
coin_change	C:FFF	C:FFF	M	C:PPP	C:PPP
cycle_detection	C:FFF	C:FFF	M	C:FFF	C:PPP
dfs	C:FFF	C:FFF	M	C:FFF	C:PPP
dijkstra	C:FFF	C:FFF	M	C:FFF	X
discrete_logarithm	C:FFF	C:PPP	C:FFF	C:PPP	C:PPP
edmond_karp	C:FFF	C:FFF	C:FFF	C:FFF	C:FFF
fast_exponential	C:FFF	C:FFF	C:PPF	C:FFF	C:PPP
gcd	C:FFF	C:FFF	C:FFF	C:FFF	C:PPP
house_robber	C:FFF	C:FFF	C:FFF	C:PPP	C:PPP
insertion_sort	C:FFF	C:PPP	C:PPP	C:FFF	C:PPP
integer_exponential	C:PPP	C:PPP	C:PPP	C:PPP	C:PPP
jump_game	C:FFF	C:FFF	M	C:FPP	C:PPP
k_smallest	C:FFF	C:FFF	C:FFF	C:FFF	C:PPP
kmp	C:FFF	C:FFF	C:FFF	C:FFF	C:PPP
knapsack_01	C:PPP	C:PPP	C:PPP	C:PPP	C:PPP
knapsack_unbounded	C:PPF	C:FFF	M	C:PPP	C:PPP
kruskal	C:FFF	X	M	C:FFF	X
lca	C:FPP	C:FPP	M	C:PPP	C:PPP
linear_search	C:PPP	C:PPP	C:FFF	C:PPP	C:PPP
linearsys_gf2	C:FFF	C:FFF	C:FFF	C:FFF	C:PPP
llrbt_delete	C:FFF	X	M	C:FFF	C:FFF
llrbt_flipcolor	C:PPP	C:PPP	C:PPP	C:PPP	C:PPP
llrbt_insert	C:FFF	X	M	C:PPP	C:FFF
llrbt_rotateleft	C:PPP	C:PPP	C:PPP	C:PPP	C:PPP
llrbt_rotateright	C:FFF	C:FFF	C:FFF	C:FFF	C:FPP
longest_common_subsequence	C:PPP	C:PPP	C:FPP	C:PPP	C:PPP
longest_increasing_subsequence	C:PPF	C:PPF	C:PPP	C:PPP	C:PPP

Continued on next page

Table 6 – continued

Task	GPT-5.5	GPT-5.6-sol	Opus think	Opus no-think	Claude Code
longest_palindrome_substring	C:FFF	C:FFF	C:FFF	C:FFF	C:FFF
matrix_multiplication	C:PPP	C:PPP	C:PPP	C:PPP	C:PPP
max_matching	C:FFF	C:FFF	C:FFP	C:FFF	C:FFF
maxheap_popmax	C:FFF	C:FFF	M	C:FFF	C:PPP
maxheap_push	C:FFF	C:FFF	C:FFF	C:FFF	C:PPP
maximum_subarray_sum	C:FFF	C:FFF	C:FFF	C:FFF	C:PPP
merge_sort	C:FFF	C:FFF	C:PPP	C:PPP	C:PPP
polymul_karatsuba	C:FFF	C:FFF	C:FFF	C:FFF	C:PPP
polymul_naive	C:PPP	C:FFP	C:FFP	C:FFF	C:FFP
prim	C:FFF	X	M	C:FFF	C:PPP
push_relabel	C:FFF	C:FFF	M	C:FFF	C:FFF
queue_dequeue	C:PPP	C:PPP	C:PPP	C:PPP	C:PPP
queue_enqueue	C:PPP	C:PPP	C:PPP	C:PPP	C:PPP
quick_sort	C:FFF	C:PPP	C:PPP	C:PPP	C:PPP
ringbuffer_dequeue	C:FFP	C:FFP	C:PPP	C:PPP	C:PPP
ringbuffer_enqueue	C:PPP	C:PPP	C:PPP	C:PPP	C:PPP
rod_cutting	C:FFP	C:PPP	C:PPP	C:PPP	C:PPP
scc_tarjan	C:FFF	C:FFF	M	C:FFF	C:FFF
segmenttree_build	C:FFF	C:FFF	C:FFF	C:FFF	–
segmenttree_modify	C:FFF	C:FFF	M	C:FFF	–
segmenttree_query	C:FFF	C:FFF	M	C:FFF	–
sieve_method	C:FFF	C:FFF	C:FFF	C:FFF	C:PPP
splaytree_splay	C:FFF	C:FFF	C:FFF	C:FFF	–
stack_pop	C:PPP	C:PPP	C:PPP	C:PPP	C:PPP
stack_push	C:PPP	C:PPP	C:PPP	C:PPP	C:PPP
string_search_naive	C:PPP	C:PPF	C:PPP	C:PPP	C:PPP
ternarysearchtree_delete	C:FFF	C:FFP	M	C:FFF	–
ternarysearchtree_insert	C:FFF	C:FFF	M	C:FFF	–
ternarysearchtree_search	C:FFF	C:FFF	M	C:FFF	–
topological_sort	C:FFF	C:FFF	M	C:PPP	C:PPP
trial_division_naive	C:FFP	C:FFF	C:PPP	C:PPP	C:PPP
trial_division_optimized	C:FFF	C:FFF	M	C:PPP	C:PPP
trie_delete	C:FFF	C:FFF	M	C:FFF	–
trie_insert	X	C:FFF	M	C:FFF	–
trie_search	X	X	M	X	–
unionfind_find	C:FFF	C:FFF	M	C:FFF	–
unionfind_linkroots	C:FFF	C:FFF	M	C:FFF	–

B.1. Semantic-Judge Disagreement Inventory

Table 7. All semantic-judge disagreements on compiler-accepted outputs.

Generation	Comparison	n	Agree	κ	Flips	Disagreement cases
GPT-5.5	5.4 T1 vs 5.4 T0	75	71	0.877	4	knapsack_unbounded (L), lca (L), rod_cutting (R), trial_division_naive (R)
GPT-5.5	5.4 T0 vs 5.6 T0	75	71	0.877	4	lca (R), longest_increasing_subsequence (L), ringbuffer_dequeue (R), rod_cutting (L)
GPT-5.6-sol	5.4 T1 vs 5.4 T0	71	69	0.939	2	lca (R), ringbuffer_dequeue (R)
GPT-5.6-sol	5.4 T0 vs 5.6 T0	71	64	0.795	7	bipartite_check (R), fast_exponential (R), llrbt_rotateright (R), longest_increasing_subsequence (L), polymul_naive (R), string_search_naive (L), ternary-searchtree_delete (R)
Opus 5 (thinking)	5.4 T1 vs 5.4 T0	49	47	0.916	2	longest_common_subsequence (R), polymul_naive (L)
Opus 5 (thinking)	5.4 T0 vs 5.6 T0	49	45	0.828	4	bipartite_check (R), fast_exponential (L), llrbt_rotateright (R), max_matching (R)
Opus 5 (no thinking)	5.4 T1 vs 5.4 T0	76	73	0.921	3	bellman_ford (R), bipartite_check (R), jump_game (R)
Opus 5 (no thinking)	5.4 T0 vs 5.6 T0	76	71	0.868	5	cycle_detection (R), fast_exponential (R), llrbt_rotateright (R), max_matching (R), polymul_naive (R)
Claude Code + Opus 5 (enhanced)	5.4 T1 vs 5.4 T0	63	61	0.881	2	bubble_sort (L), llrbt_rotateright (R)
Claude Code + Opus 5 (enhanced)	5.4 T0 vs 5.6 T0	63	61	0.871	2	bubble_sort (R), polymul_naive (R)